ARC2-Global: Tangent-Corrected Alignment for Embedding-Index Migration

Anonymous Author(s)

Anonymous Institution

Abstract. Embedding-model upgrades must preserve retrieval quality while document re-encoding proceeds under a limited budget. We propose ARC2-Global, a compact document adapter that refines centered Procrustes alignment with two ridge-fitted tangent corrections. The updates target cosine-relevant directions and collapse into one affine map, preserving native destination queries without neural mapper training. Across four development collections, two encoder migrations and three anchor-inclusive budgets, ARC2-Global achieves the highest macro nDCG@10 among six evaluated methods. Its gain over Procrustes is 0.0243, with family-adjusted interval [0.0181, 0.0308]; the second correction adds 0.0069. Matched-initialization FiQA/MiniLM controls also favor tangent targeting over chordal and Euclidean residuals, with positive corrected intervals. On prospective TREC-COVID/E5, gains over the evaluated Converter range from 0.2000 to 0.2572. A descriptive comparison across all five collections further shows E5 ARC2 at 25% coverage exceeding Converter and affine ridge at 75%. On TREC-COVID/MiniLM, 75% coverage yields 0.5195 nDCG versus native Full-Y's 0.5203, with the highest Recall@100 point mean among the six adapters. The final 384-to-768-dimensional coefficient matrix occupies 1.13 MiB, using 114.7 times fewer parameter bytes than the evaluated Converter.

Keywords: Embedding migration · Vector retrieval · Representation alignment · Mixed galleries

1 Introduction

Embedding-model upgrades create an index-maintenance problem: destination-model queries must remain effective while legacy document vectors are replaced gradually. A staged migration learns from paired embeddings, maps the remaining documents, and replaces these proxies as true destination vectors become available. We seek strong intermediate retrieval with limited destination-document exposure.

The intermediate index must rank true destination embeddings and mapped proxies together. A mapper that works well on a proxy-only gallery can change behavior after true vectors are inserted. Cosine retrieval also depends on direction, whereas coordinate regression spends capacity on radial errors that disappear after normalization. This motivates a directional document adapter that preserves compact inference and searches with the native destination query.

Document conversion and query adaptation bridge embedding spaces [18,12]; orthogonal alignment, local alignment and neural regression offer different flexibility [17,3]. We correct the prediction error that changes direction while retaining one affine document map.

ARC2-Global applies two tangent corrections to centered rectangular Procrustes. Each projects destination error onto the tangent plane at the predicted direction, rescales it to the raw prediction, and fits an affine ridge update. Both updates collapse into one affine map followed by normalization. Native destination queries search true and proxy vectors together; each encoded document replaces its proxy under the same identifier (Fig. 1).

The results identify a strong quality–exposure trade-off. On four development collections, the fixed 1,024-pair adapter leads all five baseline configurations in macro nDCG@10 across two migrations and three budgets, gaining 0.0243 over Procrustes. The second correction adds 0.0069 without adding an inference stage. On the prospective TREC-COVID collection, E5 gains over Converter and ridge persist at every budget. Across all five collections, E5 ARC2 at 25% coverage also exceeds these two alternatives at 75% in a descriptive comparison. Matched all-paid fitting strengthens MiniLM adaptation by another 0.0596 nDCG at the same paid-document budget.
Our contributions are:

- **Directional refinement with affine inference.** Two tangent-residual ridge updates target prediction direction and improve measured development retrieval while retaining one compact document map and native destination queries.
- **Strong quality at matched fitting exposure.** Across two migrations and three paid-document budgets, the four-collection development macro favors ARC2-Global over all five evaluated alternatives. Ablations test the second update and directional targeting under matched initialization; all-paid fitting tests larger supervision.
- **Useful quality at lower document exposure.** A held-collection evaluation confirms E5 gains over Converter and affine ridge at each tested budget. The descriptive 25%-versus-75% comparison exposes a useful migration operating point, while mixed-gallery diagnostics explain why adapter choice must be evaluated after native-vector insertion.

2 Related Work

Document-side translation and mixed databases. Embedding-Converter learns document-side conversion using regression together with global and local similarity losses [18]. Yang et al. study integration across embedding models through global alignment and locally fitted isometries, giving the A2M and LA2M approaches [17]. Their authors' repository also provides the SimpleLinear neural mapper used in our comparison [3]. More recently, Yang and Cao study out-of-distribution translation, multi-model mixing, and chaining, with a hierarchical mixture of experts and residual regularization along target principal

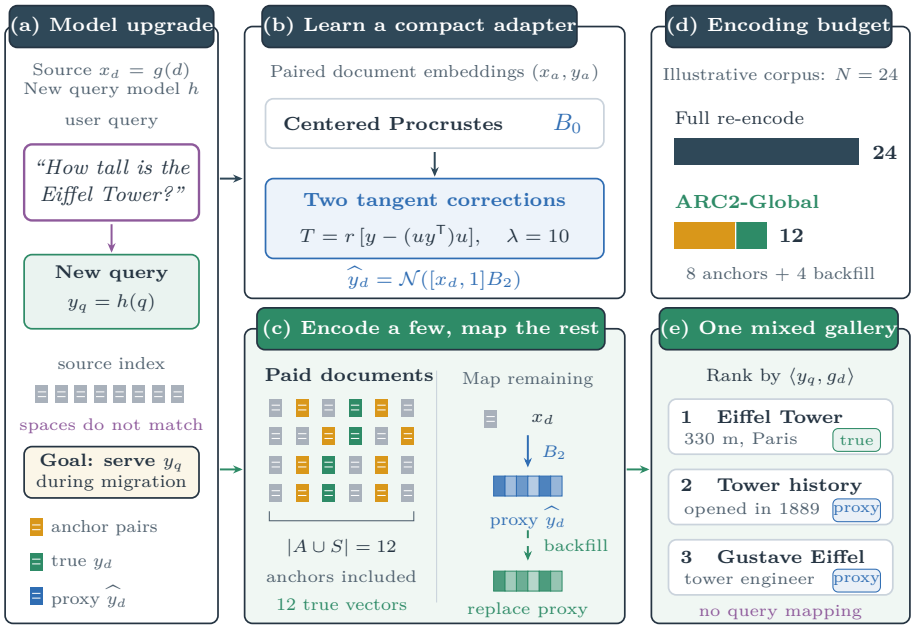


Fig. 1. A worked ARC2-Global migration. Paired document vectors train centered Procrustes and two tangent corrections, yielding one affine adapter. Unencoded documents remain searchable through proxies; backfill replaces them with native vectors in the same gallery. The 24-document corpus, eight anchors, four additional encodings and displayed hits are illustrative, not measured retrieval results.

directions [16]. Their localized expert structure addresses broader multi-model integration; our directional target uses the tangent plane at each anchor's current prediction, and both corrections collapse into one global affine map. FastFill already combines visual feature alignment with uncertainty-based partial backfilling, and shows that replacement can initially reduce retrieval performance [4]. Building on these migration settings, we study the two-correction affine computation and its relevance-judged text-retrieval behavior under anchor-inclusive encoding budgets.

Compatibility and translation direction. Backward-compatible representation learning trains a new encoder to remain usable with previously computed gallery features [10]. Drift-Adapter instead fits a post-hoc map from new-model queries to an existing legacy index, comparing Procrustes, low-rank affine, and residual neural adapters [12]. Spherical interpolation refines Procrustes-aligned queries through SLERP between old and aligned new query directions [8]; ARC2 instead fits tangent ridge updates to document anchors. ARC2-Global shares the use of paired embeddings and frozen encoders, but maps legacy document vectors into

the destination space. Its evaluation concerns mapped and native destination documents competing as re-encoding progresses.

Geometric alignment and refinement. Orthogonal Procrustes provides the classical least-squares alignment underlying our initialization [9]. Recent work connects approximate preservation of pairwise inner products to bounds on Procrustes alignment error and demonstrates interoperability across retrieval models [6]. Unpaired translation pursues a different supervision regime: vec2vec learns translation through a shared latent representation without known correspondences [5], while mini-vec2vec combines pseudo-pair construction, Procrustes estimation, and iterative matching and clustering refinements [2]. GCPA also corrects directional mismatch after Procrustes alignment [1]. It constructs a shared multi-model universe and trains a residual MLP toward consensus directions with an angular trust penalty. ARC2 instead targets paired destination documents and regularizes two affine ridge corrections around the current coefficient matrix (Eq. 8). The distinction lies in the destination target, coefficient-space prior and finite affine computation; directional post-alignment correction itself is established.

Retrieval evaluation. BEIR emphasizes retrieval across heterogeneous tasks and domains [11], while NanoBEIR provides smaller collections suitable for development comparisons [19]. Sentence-BERT, MiniLM, and E5 supply established encoder families and training approaches relevant to our model transitions [7,15,14]. Our evaluation measures judged relevance in the mixed gallery that serves destination queries, separately from vector reconstruction, native-ranking agreement or approximate nearest-neighbor overlap.

3 ARC2-Global: Angular Refinement for Mixed Galleries

ARC2-Global refines one global affine map to combine destination-space proxies with true destination vectors.

3.1 Migration Setting and Budget

Let a corpus contain N documents with stored, unit-normalized source embeddings $x_i \in \mathbb{R}^{d_s}$, written as row vectors. The destination encoder produces unit vectors $y_i \in \mathbb{R}^{d_t}$, but obtaining every y_i requires encoding the entire corpus again. We seek destination-space document representations using only a budgeted set of destination document vectors. The dimensions may differ; our evaluated encoder pairs map $d_s = 384$ to $d_t = 768$. Queries are encoded directly by the destination query encoder; the document map aligns with that existing destination query space.

Let A be the paired anchor set used for fitting and S the set selected for document replacement. The unique destination-document budget is

$$b = |A \cup S|, \qquad b = \lfloor \rho N \rfloor. \tag{1}$$

Anchors consume this budget. Our selection starts with A and fills the remaining slots from a fixed random document order, skipping duplicates, so $A \subseteq S$ and $|S| = b$. Larger budgets extend the same order and retain the same fitted anchors. The principal setting uses $m = |A| = 1{,}024$ pairs; an all-paid-pairs comparison separately fits on every paid pair. Fit uses paired document vectors without queries or relevance labels. This separates the amount of fitting supervision from the number of true vectors in the gallery. The same map serves several replacement budgets; after full replacement, scores depend entirely on the destination encoder.

3.2 Centered Rectangular Initialization

A centered Procrustes map supplies a geometric starting point before angular refinement. Stack the source and destination anchors as $X_A \in \mathbb{R}^{m \times d_s}$ and $Y_A \in \mathbb{R}^{m \times d_t}$, with row means μ_s and μ_t. Using a thin singular value decomposition, compute

$$(Y_A - \mathbf{1}\mu_t)^\top (X_A - \mathbf{1}\mu_s) = U\Sigma V^\top, \qquad W = UV^\top. \tag{2}$$

Here $W \in \mathbb{R}^{d_t \times d_s}$ is semiorthogonal and accommodates unequal dimensions. The raw prediction is $(x - \mu_s)W^\top + \mu_t$.

We expand it into affine coefficients to accumulate corrections without extra inference stages. With $p = d_s + 1$ and $\widetilde{x} = [x, 1]$, define

$$B_0 = \begin{bmatrix} W^\top \\ \mu_t - \mu_s W^\top \end{bmatrix} \in \mathbb{R}^{p \times d_t}, \qquad \widetilde{x}B_0 = (x - \mu_s)W^\top + \mu_t. \tag{3}$$

This expansion preserves raw prediction norms, without row normalization or a fitted scale; those norms enter the correction targets. The raw affine expansion uses float64 coefficients computed from the float32 Procrustes means and matrix.

3.3 Two Tangent-Residual Ridge Corrections

Write $\mathcal{N}(z) = z/\|z\|_2$ for a nonzero vector z. Cosine retrieval depends on document direction, whereas a raw vector residual also contains a radial component. We therefore project each anchor target onto the tangent plane at its current predicted direction. Let $\widetilde{X} = [X_A, \mathbf{1}]$. At update $t \in \{0, 1\}$, for anchor j set

$$\begin{aligned} &z_j^{(t)} = \widetilde{x}_j B_t, && r_j^{(t)} = \|z_j^{(t)}\|_2, \qquad u_j^{(t)} = z_j^{(t)}/r_j^{(t)}, \\ &T_j^{(t)} = r_j^{(t)} \left[y_j - \left(u_j^{(t)} y_j^\top\right) u_j^{(t)} \right]. \end{aligned} \tag{4}$$

The bracketed term is orthogonal to $u_j^{(t)}$; the norm factor restores raw prediction coordinates.

The target has a direct cosine-loss interpretation. For a nonzero row prediction z and unit destination anchor y_j, define $\ell_j(z) = 1 - y_j \mathcal{N}(z)^\top$. Differentiation gives

$$\nabla_z \ell_j(z) = -\frac{y_j - (u y_j^\top)u}{\|z\|_2}, \qquad T_j^{(t)} = -(r_j^{(t)})^2 \nabla_z \ell_j(z_j^{(t)}). \tag{5}$$

Table 1. ARC2-Global fitting and gallery construction. All document vectors supplied to fitting or replacement are unit normalized.

1 Form paired anchors X_A, Y_A and augmented design $\widetilde{X} = [X_A, \mathbf{1}]$.
2 Compute centered Procrustes and raw affine B_0 using Eqs. (2)–(3).
3 Factor $\widetilde{X}\widetilde{X}^\top + 10I_m$ once.
4 For $t = 0, 1$: compute $T^{(t)}$ from Eq. (4), solve Eq. (7), and accumulate B_{t+1}.
5 Store B_2; map documents with $\widehat{y}_i = \mathcal{N}([x_i, 1]B_2)$.
6 Replace rows in S by y_i; search with destination query vectors.

The target is a norm-scaled negative cosine gradient. Shared affine regression smooths targets across anchors, so fitted row corrections need not be tangent.

One global ridge regression shares these correction targets across anchors. Stacking their rows into $T^{(t)}$, we solve

$$\Delta_t = \underset{\Delta \in \mathbb{R}^{p \times d_t}}{\arg\min} \|\widetilde{X}\Delta - T^{(t)}\|_F^2 + \lambda\|\Delta\|_F^2, \tag{6}$$

$$= \widetilde{X}^\top(\widetilde{X}\widetilde{X}^\top + \lambda I_m)^{-1}T^{(t)}, \qquad B_{t+1} = B_t + \Delta_t. \tag{7}$$

The penalty regularizes the correction, including its intercept. Equivalently, changing variables to the updated coefficient matrix yields

$$B_{t+1} = \underset{C}{\arg\min} \|\widetilde{X}C - (\widetilde{X}B_t + T^{(t)})\|_F^2 + \lambda\|C - B_t\|_F^2. \tag{8}$$

Unlike zero-centered ridge onto Y_A, this penalty retains the current map while fitting norm-scaled cosine descent targets. We fix $\lambda = 10$ and perform exactly two updates, recomputing targets after the first; the development ablation in Sect. 5 measures the second update's contribution. A Cholesky factorization is reused across updates. The refined map is unconstrained; fitting requires paired document vectors and linear solves, with no neural mapper training.

The norm factor also gives positive global scale equivariance. In exact arithmetic with nonzero predictions, $T(cB) = cT(B)$ for a common $c > 0$. Linearity of the ridge solve therefore scales each update by c; by induction both accumulated maps scale by c, leaving normalized inference unchanged. This property concerns a common coefficient scale, not arbitrary per-document changes or guaranteed finite-step cosine improvement.

3.4 Affine Proxies and Mixed Retrieval

The two corrections collapse into the final coefficient matrix B_2. For each document, compute a proxy $\widehat{y}_i = \mathcal{N}([x_i, 1]B_2)$ and construct

$$g_i = \begin{cases} y_i, & i \in S, \\ \widehat{y}_i, & i \notin S. \end{cases} \tag{9}$$

Every anchor is replaced by its true destination vector. For a unit destination query q, rank by $qg_i^\top$. True and proxy vectors compete in one score space, so proxy approximation alone does not determine replacement effects. We use exact float32 inner products, sorting by descending score and then ascending document row index for ties; dataset-defined identical query/document IDs are excluded where specified. Table 1 summarizes the complete computation.

3.5 Computation and Storage

The final mapper has pd_t coefficients regardless of anchor count. For $384 \to 768$, this is 295,680 float32 values, or 1,182,720 bytes (1.13 MiB), including the intercept. This is coefficient storage; fitting objects, document galleries, and encoder weights are additional. Application requires one affine multiplication and normalization per document, with $O(Npd_t)$ arithmetic for a complete proxy gallery.

Fitting cost depends on the solver backend. Beyond the cross-covariance and rectangular SVD, costing $O(md_sd_t)$ and $O(d_sd_t\min(d_s,d_t))$, respectively, the reference dual implementation forms an $m\times m$ Gram matrix in $O(m^2p)$ time, factors it in $O(m^3)$, and spends $O(m^2d_t+mpd_t)$ per correction. Its Gram storage is $O(m^2)$. A separately identified primal backend solves $(\widetilde{X}^\top\widetilde{X}+\lambda I_p)\Delta_t=\widetilde{X}^\top T^{(t)}$ instead, with $O(mp^2+p^3)$ construction and factorization and $O(p^2)$ Gram storage; it supports larger paired sets. Each primal correction costs $O(mpd_t+p^2d_t)$. The equations are equivalent in exact arithmetic, while floating-point rankings can differ near ties. Updates use float64 and the final coefficient is stored as float32. Finally, $|A\cup S|$ measures logical destination-vector exposure in cached-vector evaluations. It does not measure physical encoding work, deployment latency, or end-to-end migration savings.

4 Experimental Setup

We ask three questions: does angular refinement improve fixed-supervision alignment, how does it compare with other mappers at a matched encoding budget, and which gains transfer to a prospectively held collection? All experiments use exact retrieval, isolating representation quality from approximate-index effects.

Collections and migrations. Development uses FiQA (57,638 documents, 500 development queries), NanoNQ (5,035/50), NanoHotpotQA (5,090/50), and NanoFEVER (4,996/50), drawn from BEIR and NanoBEIR [11,19]. FiQA informed method selection after its scores were inspected; all four collections are development evidence. We subsequently fixed a separate confirmation on the BEIR release of TREC-COVID [13,11]. Its input manifest contains all 171,332 documents and 50 queries. For this study, labels remained closed until every ranking and endpoint had been sealed and separately verified. This distinguishes development evidence from the prospective test. We evaluate two 384-to-768-dimensional migrations: `all-MiniLM-L6-v2` to `all-distilroberta-v1`, and `e5-small-v2` to `e5-base-v2`.[1] Model revisions and input hashes are pinned.

[1] Pinned maintainer cards: MiniLM, DistilRoBERTa, E5-small, and E5-base.

Text encoding uses normalized float32 vectors, a maximum of 128 tokens, and the E5 query/passage prefixes.

Matched supervision and exposure. The main matrix crosses four development collections, two migrations, three paid-document budgets (25/50/75%), six methods, and seeds 17, 29, 41, 53, and 67. All methods fit the same ordered 1,024 paired anchors and replace the same paid document rows. Paid sets contain all anchors and are nested, with exactly $\lfloor \rho N \rfloor$ unique target document vectors; the accounting is $|A \cup S|$, including anchors. On the Nano collections and TREC-COVID, independent seeded streams sample anchors without replacement and a document permutation; unseen permutation rows fill the budget. FiQA retains its historical anchor/order arrays. Neither fitting nor selection uses relevance labels. Paid nonanchor vectors provide replacement only in this fixed-anchor regime. A separate FiQA experiment holds the 50% paid set fixed and compares fitting on 1,024 versus all 28,819 paid pairs. Development fixed-anchor ARC2 rankings use the float64 dual backend; all-paid and TREC-COVID fits use the float64 primal backend in Sect. 3.5. The all-paid study reuses the dual fixed-anchor comparisons. Both backends implement the same two-update rule.

Matched-target controls. An additional FiQA/MiniLM study at 50% coverage holds B_0, the ordered 1,024 anchors, $\lambda = 10$, the current-map penalty and two updates fixed. It replaces the tangent target by either norm-scaled chordal error $r(y-u)$ or Euclidean error $y-z$, recomputed at each control's own state. All four arms, including unrefined B_0, run in the same Windows environment (NumPy 2.5.3, SciPy 1.18.1), using float64 dual fitting and the same float32 exact search. This is a separate development measurement.

Comparison methods. We retain the full five-baseline panel: centered rectangular Procrustes (an A2M adaptation); LA2M with 20 source-anchor clusters, local alignment for clusters with at least two anchors and global fallback for singletons; SimpleLinear; affine ridge with $\lambda = 10$ and a regularized bias; and EmbeddingConverter [17,3,18]. SimpleLinear uses pinned official code: despite its name, the locked configuration is a five-layer SELU MLP, trained for the final 50 epochs with batch size 4,028 and the MSE plus cosine loss, without qrel-based stopping. Converter is a paper-specification adaptation with a 50,000-update selection run, batch size 64, and local neighborhood size 100. Paired-vector validation MAE selects the earliest minimum at 250-update checkpoints; a fresh fit uses all fitting pairs for that selected length. Unlike the original retrieval-based validation, this adaptation uses paired-vector MAE to respect paired-document-only supervision. These are the declared implementations and training policies, rather than claims to reproduce every original-paper setting. ARC2-Global remains fixed at two tangent updates and $\lambda = 10$ throughout.

Retrieval and inference. Destination queries search a single gallery containing unit-normalized proxies and paid true vectors, using exact float32 inner products and top-100 retrieval. Score ties use original document row order; identical query/document IDs are excluded. Native Full-X and Full-Y use their own

Table 2. Development means at matched 1,024-pair supervision, equally weighting 24 conditions. Differences are ARC2 minus baseline; intervals give approximate 95% family coverage. Bold/underline mark the best/second unrounded adapter means; blue identifies ARC2.

Method	nDCG@10 ↑	R@100 ↑	Δ nDCG	Family interval
ARC2-Global	**0.5275**	<u>0.8659</u>	—	—
Procrustes	<u>0.5032</u>	**0.8661**	+0.0243	[+0.0181, +0.0308]
LA2M20	0.4234	0.7758	+0.1041	[+0.0861, +0.1229]
SimpleLinear	0.3554	0.6130	+0.1721	[+0.1499, +0.1947]
Affine ridge	0.4318	0.7441	+0.0957	[+0.0754, +0.1164]
Converter	0.3663	0.7641	+0.1612	[+0.1427, +0.1799]

query/document encoders throughout and provide reference operating points. Our primary metric is graded nDCG@10 with gain $2^{\text{rel}} - 1$; auxiliary Recall@100 measures coverage over all positive judgments, including positives outside a Nano corpus. All queries count; unjudged documents are nonrelevant. The prospective figure reports both metrics; a companion Recall table is included in the manuscript-source package.

Each locked analysis uses 20,000 paired query-bootstrap draws after averaging the five mapping seeds. Development macro effects weight the 24 collection/migration/budget conditions equally and resample queries within each collection. Within each analysis, the same sampled query indices are shared across methods, budgets and migrations for that collection. Saved 99% percentile intervals give approximate 95% coverage for the five-comparison family; the ablations have a separate five-contrast family. The all-paid experiment uses its predeclared 32-contrast family. TREC-COVID has 30 primary contrasts (five baselines, two migrations, three budgets), with approximately 99.8333% individual percentile intervals for approximate 95% family coverage. Its independent units are the 50 queries. Recall intervals are nominal 95% descriptive intervals. An interval touching zero is unresolved. The matched-target study retains its original four-contrast family, with 98.75% individual intervals; two MiniLM contrasts are observed and the two planned E5 contrasts remain pending. These intervals condition on the five observed mapping seeds and fixed collections and encoder pairs; development intervals do not adjust for method selection.[2]

5 Results

5.1 Broad Retrieval Gains with a Compact Adapter

ARC2-Global achieves the highest macro nDCG@10 among the six evaluated methods (Table 2). Its gain over Procrustes is 0.0243, with family-adjusted in-

[2] The accompanying anonymous artifact provides mapper code, baseline settings, frozen per-query results, raw ranking evidence and scripts to verify the reported numerical results.

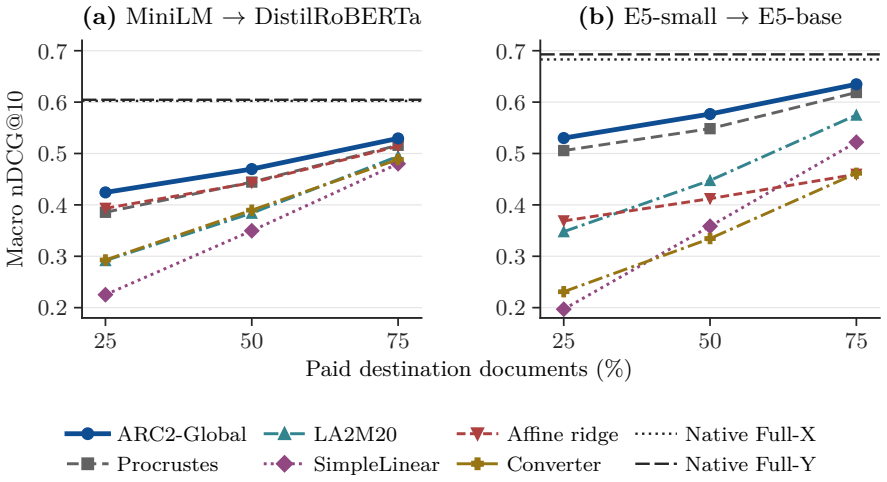


Fig. 2. Development quality at anchor-inclusive paid-document coverage. Points average five seeds and then weight the four collection means equally. All six adapters and both native references are retained. Lines connect the three measured budgets; they do not imply intermediate measurements.

terval $[0.0181, 0.0308]$. The effect extends across both migrations: descriptive pair-wise gains are 0.0257 for MiniLM and 0.0229 for E5. Each of the four collection-level macros also favors ARC2 over every baseline. At the finer condition level, ARC2 leads Procrustes in 23 of 24 point means and every other baseline in all 24. Figure 2 shows the complete budget curves.

The improvement concerns top-ranked relevance: macro Recall is 0.8659 versus Procrustes' 0.8661, with descriptive difference interval $[-0.0049, 0.0043]$. The single negative Procrustes nDCG point difference is FiQA/E5 at 75% (-0.0036). Native Full-X and Full-Y provide the full-corpus operating context.

5.2 Directional Targets Add Value without an Inference Stage

A second correction adds 0.0069 macro nDCG beyond one step, with family interval $[0.0038, 0.0100]$. The point gain appears in 22 of 24 conditions and in each of the five seed-level macros. Both updates retain one affine inference stage. The complete initialized refinement gains 0.0957 over affine ridge (Fig. 3a), jointly evaluating the target and alignment prior.

The penalty controls favor moderate regularization. Changing λ from 10 to 1 gives $+0.0032$, with an interval crossing zero; changing it to 100 gives -0.0093, with a wholly negative interval. The evaluated method keeps $\lambda = 10$ and exactly two corrections.

Tangent targeting improves matched residual controls. Holding initialization and supervision fixed isolates an additional benefit on FiQA/MiniLM (Fig. 3b).

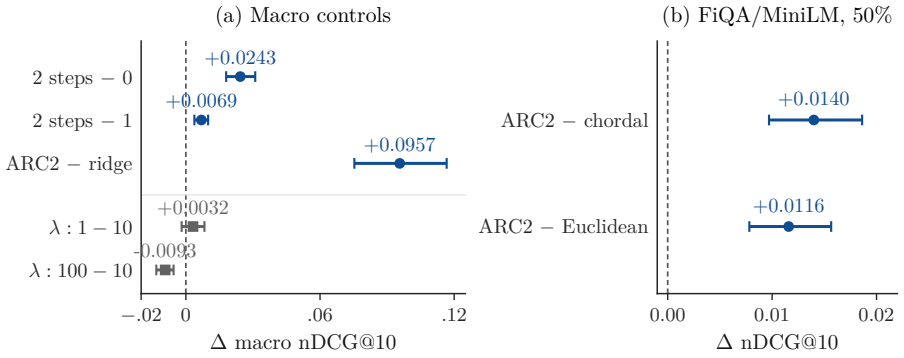


Fig. 3. Development mechanism controls. (a) Equal-condition macro effects, with saved 99% intervals for five contrasts. (b) FiQA/MiniLM at 50% coverage: matched-initialization target controls, with 98.75% intervals retaining the original four-contrast family. Panels use separate scales; all intervals are shown in full.

ARC2 reaches 0.3075 nDCG, versus 0.2935 for the norm-scaled chordal control, 0.2959 for Euclidean residual refinement and 0.2858 for unrefined Procrustes. The tangent gains are 0.0140 ($[0.0097, 0.0186]$) and 0.0116 ($[0.0078, 0.0157]$), respectively, with both differences positive in all five fixed seed means. The chordal comparison supports removing the radial target component in this development configuration. It shares the norm-scaling rule, without imposing identical later states or equal target norms.

5.3 Prospective Transfer: Strong E5 Gains and MiniLM Coverage

On the held collection, E5 ARC2 exceeds the evaluated Converter and affine ridge at **all three migration budgets**. At 25% coverage, the Converter gain is 0.2572, with family interval $[0.1757, 0.3376]$. At 50%, ARC2 reaches 0.5904 nDCG versus Converter's 0.3590 and ridge's 0.3281 (Table 3). Across both migrations, the 30-contrast family supports five improvements over Converter, four over ridge and two over SimpleLinear. The table marks all eleven supported comparisons.

MiniLM provides a useful partial-migration operating point: at 75% coverage, ARC2 reaches 0.5195 nDCG versus native Full-Y's 0.5203, while one quarter of the gallery still uses proxies. Its Recall@100 point mean is the highest among the six adapters at each budget, reaching 0.0998 at 75% versus native Full-Y's 0.1052. E5 reaches 0.0997 Recall at 75%, versus Full-Y's 0.1162 (Fig. 4). These auxiliary results distinguish coverage at 100 from relevance near the top; proximity in nDCG is not an equivalence test.

Geometric alignment remains competitive: before rounding, Procrustes has higher nDCG point means in all six cells, although its MiniLM/75% difference is below 0.0001. All Procrustes and LA2M nDCG family intervals cross zero.

Table 3. Prospective TREC-COVID nDCG@10 (higher is better), averaged over five seeds. †: positive ARC2-minus-baseline interval in the 30-contrast family. Unmarked baseline cells are unresolved. Bold/underline rank the six adapters before rounding; native references are separate.

Method	MiniLM → DistilRoBERTa			E5-small → E5-base		
	25%	50%	75%	25%	50%	75%
ARC2-Global	0.4619	0.5000	<u>0.5195</u>	<u>0.5628</u>	<u>0.5904</u>	0.6098
Procrustes	<u>0.4620</u>	0.5046	**0.5195**	**0.5660**	**0.6067**	**0.6198**
LA2M20	**0.4634**	<u>0.5073</u>	0.5186	0.5206	0.5898	<u>0.6188</u>
SimpleLinear	0.4610	**0.5076**	0.5193	0.4299†	0.5287†	0.5785
Affine ridge	0.4322	0.4854	0.5012†	0.2986†	0.3281†	0.3278†
Converter	0.4224†	0.4780†	0.5033	0.3056†	0.3590†	0.4098†
Native Full-X	0.4935 (native)			0.6805 (native)		
Native Full-Y	0.5203 (native)			0.6421 (native)		

Table 4. E5 nDCG: ARC2 at 25% exceeds Converter and ridge at 75% on all five collections. Blue identifies ARC2. Proc.: Procrustes; SL: SimpleLinear; native references show full-corpus quality. TREC-COVID is prospective; other rows are development. Cross-budget point comparisons are post-hoc descriptive.

Collection	25% coverage			75% coverage			Native	
	ARC2	Proc.	LA2M20	SL	Ridge	Conv.	Full-X	Full-Y
FiQA	0.2731	0.2718	0.1416	0.2781	0.1610	0.2095	0.3465	0.3891
NanoNQ	0.5071	0.4852	0.3640	0.4839	0.4428	0.4021	0.6595	0.6669
NanoHotpotQA	0.5845	0.5512	0.3869	0.6174	0.5053	0.5016	0.8278	0.8193
NanoFEVER	0.7555	0.7147	0.4991	0.7087	0.7290	0.7326	0.8984	0.8963
TREC-COVID	0.5628	0.5660	0.5206	0.5785	0.3278	0.4098	0.6805	0.6421

The confirmed advantages therefore concern the evaluated Converter, ridge and SimpleLinear configurations.

5.4 Higher E5 Quality at Lower Document Exposure

A useful budget pattern extends across all five collections (Table 4): E5 ARC2 at 25% coverage exceeds both Converter and affine ridge at 75%. On TREC-COVID, the respective nDCG means are 0.5628, 0.4098 and 0.3278. The advantage is visible even after the two comparators receive a much larger set of true destination vectors.

This post-hoc descriptive comparison lies outside the matched-budget inferential family. The table retains geometric alternatives, SimpleLinear and native references. Physical savings and minimum equal-quality budgets require separate measurements.

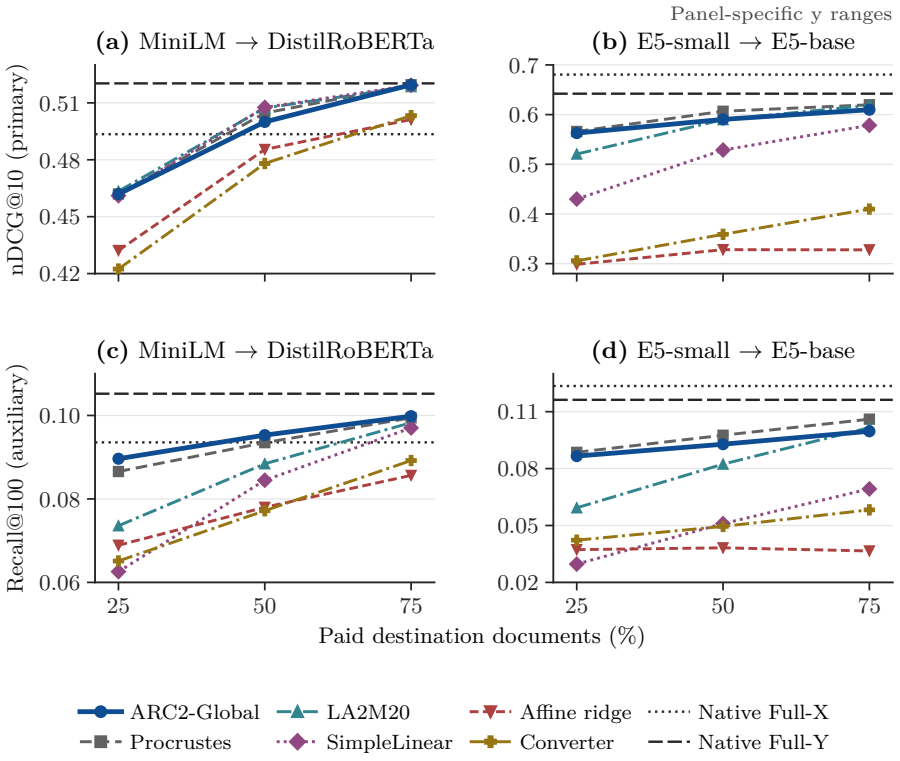


Fig. 4. Prospective TREC-COVID at fixed 1,024-anchor supervision: (a,b) primary nDCG@10; (c,d) auxiliary Recall@100. Points average five seeds. All six adapters and native references are shown, with explicitly panel-specific vertical ranges. Recall curves are descriptive.

5.5 Paid Pairs Strengthen MiniLM Adaptation

Using already-paid pairs improves MiniLM adaptation without increasing the 50% document budget. ARC2 rises from 0.3075 to 0.3671 nDCG, a gain of 0.0596 with family interval $[0.0393, 0.0806]$. In matched all-paid comparisons it exceeds Procrustes, LA2M, SimpleLinear and Converter; ridge is unresolved. Table 5 makes these family decisions visible.

E5 changes the method ordering: ARC2 trails all-paid Procrustes and LA2M, exceeds SimpleLinear and ridge, and is unresolved against Converter. Its fixed-to-all-paid change, 0.3116 to 0.2919, is unresolved. Fixed-anchor MiniLM ARC2 also trails every all-paid baseline in the cross-regime family comparisons. All-paid fitting retains $\lambda = 10$ in the summed objective; it tests the fixed procedure at larger supervision, rather than isolating data quantity at a fixed effective penalty.

Table 5. FiQA at 50% coverage. Fixed/All fit 1,024/28,819 pairs; N/R denote nDCG@10/Recall@100. On All-N baseline cells, †/‡ mark positive/negative ARC2 differences in the 32-contrast family; unmarked comparisons are unresolved. Bold/underline rank matched adapter columns; native rows are references.

Method	MiniLM → DistilRoBERTa			E5-small → E5-base		
	Fixed N	All N	All R	Fixed N	All N	All R
ARC2-Global	**0.3075**	<u>0.3671</u>	<u>0.6922</u>	**0.3116**	0.2919	0.5595
Procrustes	0.2858	0.3334†	0.6905	<u>0.3100</u>	**0.3362**‡	**0.6534**
LA2M20	0.2377	0.3389†	**0.6950**	0.2300	<u>0.3293</u>‡	<u>0.6473</u>
SimpleLinear	0.2360	0.3323†	0.6778	0.1942	0.1679†	0.3618
Affine ridge	<u>0.2890</u>	**0.3673**	0.6800	0.1479	0.2270†	0.4245
Converter	0.2496	0.3258†	0.6633	0.1397	0.2625	0.5789
Native Full-X	0.3900 (native)		0.7002	0.3465 (native)		0.6453
Native Full-Y	0.4047 (native)		0.7039	0.3891 (native)		0.6968

Evaluate the gallery that will be searched. With each fitted map fixed, true-vector replacement raises MiniLM all-paid nDCG from 0.3510 to 0.3671, but lowers E5 from 0.3149 to 0.2919. For E5, Recall falls from 0.6542 to 0.5595 and the paid-document share of top-10 results falls from 0.4916 to 0.0647, despite 50% paid coverage. These measured replacement effects make mixed-gallery relevance the appropriate selection criterion: proxy-only quality can favor a different operating point from the gallery containing native and mapped vectors.

6 Practical Value of Directional Adaptation

Compact inference with stronger development ranking. The 295,680 float32 coefficients occupy 1.13 MiB, or 114.7 times fewer parameter bytes than the evaluated Converter. Both corrections are absorbed into this matrix. Affine ridge has the same coefficient count; refinement provides the measured quality benefit, while affine mapping provides compactness.

When the adapter is useful. ARC2 targets required destination-model queries during staged document replacement. E5 demonstrates useful quality at limited exposure; MiniLM benefits from fitting all paid pairs without increasing document coverage. Procrustes and LA2M remain competitive on the held collection, so deployment choice should follow encoder-specific mixed-gallery validation. When source queries remain acceptable, Full-X is also viable: all 24 partial-development five-seed ARC2 condition means are below it.

Efficiency in the measured units. Destination-vector exposure, coefficient storage and the absence of neural mapper training describe the demonstrated resource advantages. Physical encoding time, index resources and serving SLOs require deployment measurements.

7 Conclusion

ARC2-Global turns two tangent ridge corrections into one compact affine adapter. Matched-initialization MiniLM controls support its directional target. It leads five alternatives in development macro nDCG, transfers substantial E5 gains over Converter and ridge to TREC-COVID, and exposes useful 25%-versus-75% quality–coverage operating points. Stronger all-paid MiniLM adaptation and measured replacement effects further show why migration should be evaluated in the mixed gallery that serves destination queries. Angular refinement offers a practical route to budgeted index migration.

References

1. Achara, A., Gaintseva, T., Mahaut, M., Chakraborty, P., Johansson, V.S., Barsbey, M., Rodolà, E., Crisostomi, D.: Multi-way representation alignment. In: Proceedings of the 43rd International Conference on Machine Learning. Proceedings of Machine Learning Research, vol. 306, pp. 310–331. PMLR (2026), https://proceedings.mlr.press/v306/achara26a.html
2. Dar, G.: mini-vec2vec: Scaling universal geometry alignment with linear transformations (2026), https://arxiv.org/abs/2510.02348
3. DBgroup Edinburgh: Embedding translation: Reference implementation. Software repository (2026), https://github.com/DBgroup-Edinburgh/embedding-translation, revision 5e62baf15442; SimpleLinear implementation
4. Jaeckle, F., Faghri, F., Farhadi, A., Tuzel, O., Pouransari, H.: Fastfill: Efficient compatible model update. In: International Conference on Learning Representations (2023)
5. Jha, R., Zhang, C., Shmatikov, V., Morris, J.: Harnessing the universal geometry of embeddings. In: Advances in Neural Information Processing Systems. vol. 38, Main Conference, pp. 45963–45987 (2025). https://doi.org/10.52202/085713-1530
6. Maystre, L., Gonzalez, A.O., Park, C., Dolga, R., Berariu, T., Zhao, Y., Ciosek, K.: When embedding models meet: Procrustes bounds and applications (2025), https://arxiv.org/abs/2510.13406
7. Reimers, N., Gurevych, I.: Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP). pp. 3982–3992 (2019). https://doi.org/10.18653/v1/D19-1410
8. Ricci, S., Biondi, N., Pernici, F.: Spherical interpolation for backward-compatible multimodal representations (2026), https://arxiv.org/abs/2609.39836
9. Schönemann, P.H.: A generalized solution of the orthogonal procrustes problem. Psychometrika **31**(1), 1–10 (1966). https://doi.org/10.1007/bf02289451
10. Shen, Y., Xiong, Y., Xia, W., Soatto, S.: Towards backward-compatible representation learning. In: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) (2020)
11. Thakur, N., Reimers, N., Rücklé, A., Srivastava, A., Gurevych, I.: Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. In: Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks. vol. 1 (2021)

12. Vejendla, H.: Drift-adapter: A practical approach to near zero-downtime embedding model upgrades in vector databases. In: Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing. pp. 15938–15949 (2025). https://doi.org/10.18653/v1/2025.emnlp-main.805
13. Voorhees, E., Alam, T., Bedrick, S., Demner-Fushman, D., Hersh, W.R., Lo, K., Roberts, K., Soboroff, I., Wang, L.L.: Trec-covid: constructing a pandemic information retrieval test collection. ACM SIGIR Forum **54**(1), 1–12 (2020). https://doi.org/10.1145/3451964.3451965
14. Wang, L., Yang, N., Huang, X., Jiao, B., Yang, L., Jiang, D., Majumder, R., Wei, F.: Text embeddings by weakly-supervised contrastive pre-training (2024), https://arxiv.org/abs/2212.03533
15. Wang, W., Wei, F., Dong, L., Bao, H., Yang, N., Zhou, M.: Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. In: Advances in Neural Information Processing Systems. vol. 33, pp. 5776–5788 (2020)
16. Yang, B., Cao, Y.: Generalizable and composable multi-model embedding translation. In: Proceedings of the 43rd International Conference on Machine Learning. vol. 306, pp. 144275–144302 (2026), https://proceedings.mlr.press/v306/yang26e.html
17. Yang, B., Cao, Y., Ren, Y.: Integrating vector databases across embedding models. Proceedings of the ACM on Management of Data **3**(6), 1–28 (2025). https://doi.org/10.1145/3769803
18. Yoon, J., Arik, S.O.: Embedding-converter: A unified framework for cross-model embedding transformation. In: Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). pp. 25464–25482 (2025). https://doi.org/10.18653/v1/2025.acl-long.1237
19. Zeta Alpha: NanoBEIR. Hugging Face dataset collection (2024), https://huggingface.co/collections/zeta-alpha-ai/nanobeir, accessed 1 November 2026